

Hands-On Experience: A Pedagogical RISC-V Simulator

A Thesis
Presented to
The Division of Mathematical and Natural Sciences
Reed College

In Partial Fulfillment
of the Requirements for the Degree
Bachelor of Arts

Sebastian Andrada Ottonello

May 2024

Approved for the Division
(Computer Science)

Charles McGuffey

Acknowledgements

I would like to thank Charlie for advising me on this thesis, and for teaching the Computer Systems course that inspired it; and I would like to thank everyone involved with the RIPES simulator and the RISC-V computer architecture for similar reasons. I would like to thank the friends I've made at Reed, and those I've kept outside of it, for the support and kindness they continue to show me. Out of these, I thank Tucker and Alex in particular: The former's feedback was invaluable in my final days of editing, and I don't think I could have seen this project to completion without the latter's incredible patience and expertise that she so readily shared with me. I would like to thank my partner, Dunham, for always encouraging me onwards and for believing in me even when I myself did not. I would also like to thank my parents, Gisela and Gabriel, for their unconditional love and their insistence that I follow my dreams.

List of Abbreviations

ALU	Arithmetic Logic Unit
CISC	Complex Instruction Set Computer
CPU	Central Processing Unit
GUI	Graphical User Interface
ISA	Instruction Set Architecture
RISC	Reduced Instruction Set Computer
RISC-V	The RISC-V Instruction Set Architecture
RV32I	RISC-V 32-bit Base Integer Instruction Set

Table of Contents

Introduction	1
Chapter 1: Background	3
1.1 An abstract view of how a CPU functions	3
1.2 What is an ISA?	4
1.3 Why RISC-V?	5
1.4 RISC-V Instructions	6
1.5 RISC-V Hardware Specifications	7
1.6 A more detailed look at how a CPU functions	9
Chapter 2: CPU Design	11
2.1 Pipelining	11
2.1.1 The Backbone of Pipelining: Latches	13
2.1.2 Data Hazards	15
2.1.3 Stalling	17
2.1.4 Forwarding	19
2.1.5 Control Hazards	21
2.2 Final CPU Design	22
Chapter 3: Simulator Design	23
3.0.1 First Draft	23
3.0.2 Final design	24
3.0.3 Structs and Sub-Structs	26
3.0.4 Update Functions	29
3.0.5 Stalling	29
3.1 CPU-Simulator Features	31
3.1.1 Implementation of User Input and Rewinding	31
3.1.2 Graphical User Interface	32
Conclusion	35
4.1 Future Extensions	35
4.1.1 Data Memory	35
4.1.2 Out-of-Order Execution and Register Renaming	36
4.2 Takeaways	38
4.3 Closing Thoughts	40

References	41
----------------------	----

List of Figures

1.1	Basic CPU Design.	9
2.1	Program execution over time.	12
2.2	Program execution over time, with pipelining.	12
2.3	CPU with Pipelining Latches.	14
2.4	Program execution without pipelining.	15
2.5	Program execution with pipelining.	16
2.6	Program execution with no Stalling. Data Hazard occurs.	17
2.7	Program execution with Stalling. No Data Hazard.	17
2.8	CPU Design with stall control.	18
2.9	Program execution without Stalling	19
2.10	The three major types of forwarding.	20
2.11	Final CPU Design.	22
3.1	Color-coded CPU design. Red is for data stored in the State struct, Blue is for data in the Logic struct. Multiplexors are definitionally stateless, and are all in the Logic struct.	25
3.2	The Simulator's command-line-based GUI.	32

List of Listings

3.1	Barebones layout of <i>components.rs</i>	25
3.2	Snippet of <i>components.rs</i> , holds main 2 structs	26
3.3	<i>decode.rs</i> , the file that holds ID-stage structs	27

Abstract

Computer Systems is a critical field of computer science, but it can also have a rather harsh introduction for undergraduate students. Many of them go into introductory Computer Systems courses with little to no experience in lower-level programming languages, never mind any assembly; and while one can learn well enough how a computer microprocessor functions over the first few weeks, it may take quite some time for them to grow familiar with the topic. This is especially true for trying to learn how information flows through a chip to make the execution of computer programs possible.

In order to help such students gain a more intuitive understanding of CPUs as they are taught in introductory Computer Systems courses, I have developed a CPU simulator made explicitly for educational purposes. The simulator focuses on simplicity, ease-of-use, and transparency.

Simple CPU simulators, such as SPIM MIPS, are already in use throughout college classrooms for this sort of purpose, but they are focused solely on displaying the results of a program, and keep the simulated CPU's state opaque. Therefore, the primary goal of my simulator is not to obtain the results of running assembly programs, but rather to cleanly display its own internal state while doing so and be easy to interact with. It gives potential users the opportunity to not only write and run their own low-level assembly code programs, but to then also watch them in action; or, if a user is interested in a particular feature of CPUs, they can see how that feature is used in a variety of different scenarios. The simulator is not meant to teach how a CPU works on its own, but rather to function as an interactive tool that can supplement existing instruction.

Introduction

Computer Systems is the study of the physical, metal-and-silicon devices that make up a functioning computer, and how they operate at their lowest level. These devices are the backbone of just about any work done on a modern computer, and the field is therefore well worth exploring.

Of the countless pieces of hardware to study and delve into within this field, the CPU (Central Processing Unit) is the most essential one. It is the beating heart of any computer, the piece that actually carries out the programs and computations that it was built for. It comes as no surprise, then, that most college-level Computer Systems courses begin by taking a close look at the inner workings of a simple CPU, so that they can then begin to understand how the other components of a computer support its functionality, and how it can produce computation out of just an assembly program, wires, and logic gates. Learning of this is all well and good, but getting hands-on experience with the topic is harder than one might expect. The powerful CPUs in the personal computers of students run on are all closed-source designs that obfuscate how they carry out their instructions; and even if they were transparent, the sheer breadth and depth of the optimizations they employ to get the performance they do would make their design completely indecipherable to a novice.

Therefore, there is a clear need for an educational CPU simulator that can demonstrate how modern CPUs function in an observable, approachable, and interactive manner. This project aims to develop such a program in the Rust programming language, one that is capable of running on modern computers. This program simulates a simple RISC-V CPU at the sub-component level, running assembly code and then transparently displaying how the code and computed information traverse the inside of the processor. Here, "sub-component level" means that the CPU is arbitrarily divided up into a model of individual pieces with specific purposes, which are wired together at specific spots to form the CPU. This model is commonly used to introduce students to CPU design, so it was the obvious choice for this simulator.

While simple, the simulated CPU implements a 5-stage pipeline design similar to the ones that are mandatory for modern CPUs. The project also seeks to include some user-facing features, like a graphical interface and the ability to step through and rewind through the execution of a program, to make the simulator more intuitive to use.

The results are promising. The finished simulator supports the 32-bit RISC-V Base Integer specifications, and features a GUI, support for user-inputted programs, and rewinding capabilities for its ease of use as well. Future work could be useful

towards implementing certain useful features that were deemed out of the scope of this thesis, such as the ability to edit CPU values mid-program or a higher-fidelity simulation of the Data Memory.

Chapter 1

Background

1.1 An abstract view of how a CPU functions

In order to build and understand the design of a CPU, we must first understand what it, in general terms, does. A CPU carries out a program consisting of instructions from its ISA (Instruction Set Architecture); more information on ISAs will be provided in the next section. However, what exactly do those instructions demand the CPU do, and what does the CPU have access to to carry it out?

We will go over the exact design of the model of CPU shown here later, with all of its intricacies and small parts. But, in broad terms, the CPU has access to: an **Instruction Memory**, where the program to be executed is stored instruction-by-instruction. A **Decoder** that reads this raw instruction, made of bits, and gleans the important details from it, like what instruction it is, what the inputs and outputs are, or other instruction-specific data. A **Register Memory**, which holds a small amount of numbered, general-purpose registers that store information; think of them as set mini-variables built directly into the CPU that it has quick and easy access to. A **Data Memory**, which is external and much larger than the previous two, in which to store larger values; and finally an **ALU** (Arithmetic Logic Unit) which is simply capable of performing arithmetic operations on the values it is given.

This explanation leaves out many crucial components and smaller parts that make the whole thing possible, but they're enough to understand the CPU's general plan, detailed here below. Note that this plan, and the designs described later, do not accurately describe all CPUs; rather, it is a simple design that, despite that, is fully equipped to execute base RISC-V programs.

1. Read the first instruction from the Instruction Memory. (**IF: Instruction Fetch**)
2. Decode the fetched instruction into useful information: The operation to be performed, the registers to be used, etc. (**ID: Instruction Decode**)
3. Read the values of whichever registers the instruction calls for from the Register Memory. (**ID: Instruction Decode again.**)

4. Perform whichever arithmetic operation the instruction calls for, using values from the registers or from the instruction itself as input. (**EX: Execute**)
5. If the instruction calls for a value to be read from or stored to the Data Memory, do so. (**MEM: Memory**)
6. Write back the end-result value of the instruction into the destination register it specified. (**WB: WriteBack**)

Shockingly, these six repeating steps are enough to accomplish everything that base version of RISC-V calls for! From there, the instructions in the RISC-V ISA can be used to build any program, from the firmware of the simplest embedded device to a modern operating system. Exploring how exactly that happens is way beyond the scope of this thesis, but it is important to know that it can be done.

Note that the dividing of the process of executing one instruction into six steps is entirely arbitrary. You may have noticed that the steps have been given titles in bold, but steps 2 and 3 are both lumped into the same "ID" name; These names are called "stages", and my 6 steps and the 5 stages are both arbitrary. You could subdivide this exact same process into less or more parts depending on how deep you look. However, this arbitrary model of "stages" is going to be incredibly helpful when implementing the pipeline design of this simulator, which will be described later.

This simplified explanation ignores many practical concerns, but one is particularly important to the CPU's design, relating to steps 1 and 2. In the modern day, we want multiple CPUs of different designs to all be capable of running the same machine code. How does the CPU know which set of bits correspond to which instruction? How does it know which operations its designers should ensure it can compute?

1.2 What is an ISA?

These questions all have the same answer; and so does an entirely different question, one that you might've asked yourself back in the introduction: what does it even mean that a CPU is *RISC-V*?

The answer is, ISAs. ISAs are how CPUs know how to read machine code even when it is not tailored to their specific model; and RISC-V is one of many ISAs, belonging to the RISC family of design. So, what are they?

Starting from the top, ISA stands for Instruction Set Architecture; you may have also heard it referred to as just a CPU's "architecture", or you might've heard of popular ISAs like x86 and ARM. We will go into the details later, but for now an ISA is simply a list of machine-code instructions that a CPU can carry out: These instructions are things like basic arithmetic operations (addition, subtraction, etc), calls to load or store data from the memory, or even jump instructions to alter the flow of the program. The ISA also specifies exactly how to decode any machine-code instruction, stored as a set of bits. Ultimately, a program is just a list of these instructions, one after the other, that the CPU executes in order... or, well, that it pretends to execute in order. More on that later.

The ISA also contains certain hardware specifications that the CPU should be assumed to have: Things like the amount of registers available and how large a value they can store, or how the program counter ought to be measured, or even certain demands to make life easier for programmers. For example, the MIPS ISA demands that the 0th register, "Register r0", be hardcoded to the value zero at all times, so that programmers have easy access to the value and to give the option of a null register to point an instruction's destination at when its result is unwanted.¹ While centered around the instructions, these additional demands are a glimpse into how important and widespread the ISA is in a CPU's design; RISC-V's specifications will be described later.

A CPU is considered to use a certain ISA if it supports and has implemented all of its instructions and hardware requirements. ISAs were created for the goal of hardware compatibility: as long as two different CPUs made by competing manufacturers both use the same ISA, any low-level program will run on both processors because they both guarantee that they'll support the same instruction set, no matter how different their underlying design is.

Ultimately, an ISA is a standard for how a CPU should process any given line of code from a low-level program. This close to the hardware, the line is encoded as an integer in binary; and, by referencing the ISA, it can be translated into an actionable operation that the CPU should execute when it reaches that part of the program.

1.3 Why RISC-V?

Though the ISA doesn't prescribe an exact design for how its demands should be met (that's called the *implementation*, and is the domain of the CPU's manufacturer), it so heavily shapes the CPU's capabilities that it informs major parts of a CPU's design. Choosing which ISA to implement is the most important decision one can make when designing a CPU, whether it's real or virtual. As discussed in the introduction, there are multiple ISAs to choose from, and many are more popular ones than RISC-V. So, what makes competing ISAs so different from one another, and why would I choose RISC-V for this project instead of any of the more popular options?

RISC-V is a modern ISA, albeit one that has not yet experienced widespread success; while it very rapidly growing in use, its two contemporaries, x86 and ARM, almost entirely dominate the Personal Computer CPU and Mobile CPU market for the time being. And yet, RISC-V has two major advantages over them that make it the ideal choice for this simulator.

The first is that RISC-V is fully open source. Unlike its two proprietary cousins, the full listing of all the RISC-V instructions is freely available, with details on what behavior is expected of each instruction and exactly how to glean an instruction's parameters from its 32-bit integer representation. Working with open-source software and specifications benefits everyone, not just its creators. It's also a better choice practically: an open-source ISA is easier to implement and easier for students to study.

¹Waterman et al. (2017)

The second difference is between two opposing schools of design for ISAs: *Reduced Instruction Set Computer* (RISC) and *Complex Instruction Set Computer* (CISC). The difference ultimately comes down to how many different instructions the ISA lists out: The idea of a CISC ISA is that by providing a very large variety of different instructions, most individual actions you'd want a program to do can be performed with just one instruction each. By contrast, a RISC ISA aims to keep its total number of instructions low. x86 and ARM are both CISC ISAs, while RISC-V is a RISC ISA (it's even in the name).

While this means that code for RISC ISAs needs more total instructions to accomplish the same task, the idea is that keeping the demands simpler would allow RISC processors to have faster and more efficient designs that offset this cost.² This also coincidentally makes RISC processors a much better example for students to study; and, despite the simple design, RISC-V is still a modern, relevant architecture that's growing in widespread use.

1.4 RISC-V Instructions

This information on CPUs and ISAs is also enough to get a rough idea of how any given RISC-V instruction is carried out. The full list of instructions can be found on the official *RISC-V Instruction Set Manual* that is freely available online, or on a number of more user-friendly guides like msyksphinz's *riscv-isa-pages*.^{3 4} A list of all the instructions in the most basic version of RISC-V, the RV32I Base Integer Instruction Set, can be found in the first chapter of the aforementioned RISC-V manual, but these instructions are broadly divided into the following four categories by the *RISC-V Instruction Set Manual*:

Register-Register Computation: These instructions take two registers as input, perform some sort of operation on their two values (an arithmetic operation like addition, or a logical operation like a bitwise AND or NOT), and store the result in a third register.

Register-Immediate Computation: These are the same as RR (Register-Register) instructions, but take as input one register and one hard-coded value that's written directly into the program, called an Immediate. It's the difference between " $z = x + y$ " and " $z = x + 3$ ".

Load/Store: The purpose of these instructions are simply to interact with the Data Memory, either storing a value to a specific address or reading a value from it.

Control Transfer: These are instructions that call for the CPU to move on to a place in the code that is NOT the next instruction, potentially. These are what make the If-Then statements and Loops of high-level programming languages possible. These work because there is a simple **Program Counter** component, a PC, attached to the Instruction Memory that simply keeps track of which instruction should be performed next.

²Denning (1993)

³Waterman et al. (2017)

⁴msykpshinz (2019)

Each of these instructions in these categories, conveniently, move through the 5-stage process in roughly the same way based on their category:

RR Computation: Run through the steps as usual, feeding the value of its two input registers as the input for the ALU. ignore step 4, MEM, since Computational Instructions do not interact with the Data Memory.

RI Computation: Same as RR; but only one register is read at step 2. The other input value is the immediate, which the Decoder reads from the instruction itself.

Load: Load Instructions have a Register and an Immediate as input, representing the address in data memory to read from, as well as a destination register. They run through all of the steps; in the EX stage, the register value and immediate are added together to get the final address, which is then fed to the Data Memory in the MEM stage. Finally, the data loaded from Data Memory is stored in the destination register.

Store: Store Instructions are similar to Load, but have two registers and an Immediate, with no destination register. The first register + Immediate are the address as usual, but the second register holds the value that is to be stored into memory; and is sent directly to the Data Memory through its own special wire.

Control Transfer: Control Transfer Instructions have an Immediate that describes where the program should jump to, and possibly some registers. They run normally until the EX stage: Here, rather than operate on registers, the ALU is given the Immediate and the current value of the PC, which are added together to determine the Jump location. Some additional computation may be performed in this stage, such as a comparison between registers, and for some instructions the jump destination is written to an output register, but the CPU, at this point, goes back to the IF stage while editing the value of the Program Counter to complete the jump.

Do note that these categories are loose and arbitrary, to an extent. The RISC-V Instruction Manual itself has six "Types" it divides its instructions into, but these have less to do with what the instruction accomplishes and more with the details of how the CPU should read its immediates bits. The categories described above help give a broad overlook of what sort of instructions the RISC-V ISA expects CPU implementations to handle.

There are also a handful extra instructions in the RV32I that only make sense in the context of an operating system that a program would be running on top of, used for things like system calls, debugging, and accessing privileged hardware timers. Since this simulator supposedly runs its programs bare-metal with no OS beneath it, these instructions are useless in the context of the project and will not be supported.

1.5 RISC-V Hardware Specifications

Despite what was stated earlier, the ISA does also makes certain demands, or specifications, for the hardware of the CPUs that implement it. Generally, the hardware side of things is the CPU's job, not the ISA; but in order to be compatible with every CPU running under it, an ISA needs to have certain guarantees. What is a RISC-V program tries to write a value to register "r19", but the CPU running it has only 16

registers?

RISC-V has two base variations to pick from, The *RV32I Base Integer Instruction Set* and the *RV64I Base Integer Instruction Set*. It also has many optional extensions, which are part of why RISC-V can have a simple core while remaining viable for commercial applications. However, it is the Base Integer Instruction Sets, referred to as RV32I and RV64I respectively from now on, that determines hardware specification, so that is what we will look at now.

RV32I calls for a total of 32 registers. Registers r1 through r31 are general-purpose registers, each capable of holding a 32-bit-wide integer value. r0 is a special register hardwired to the constant value 0, which is useful for many situations.

It specifies that each Instruction should be 32-bits in length, though it doesn't specify how many Instructions the Instruction Memory should be able to hold. And, while the ISA doesn't make any demands of the Data Memory, and actually warns that the environment a program is executed in (like, say, an Operating System) gets to decide which parts of the Data Memory it can access, the amount it can access is upper-bound by the addresses that the RV32I Load and Store instructions can encode in their 32 bits ⁵.

RV64I is rather similar, with the difference that the 32 registers have a width of 64 bits each, and that Instructions themselves are 64-bits long ⁶.

The RV32I is simpler, and easier to understand at a glance, making it a better choice for an educational simulator.

⁵Waterman et al. (2017)

⁶Waterman et al. (2017)

1.6 A more detailed look at how a CPU functions

This information is enough to build a full hardware model of the CPU that the simulator will use, at least if pipelining is ignored for now. The following is a visual representation of this model:

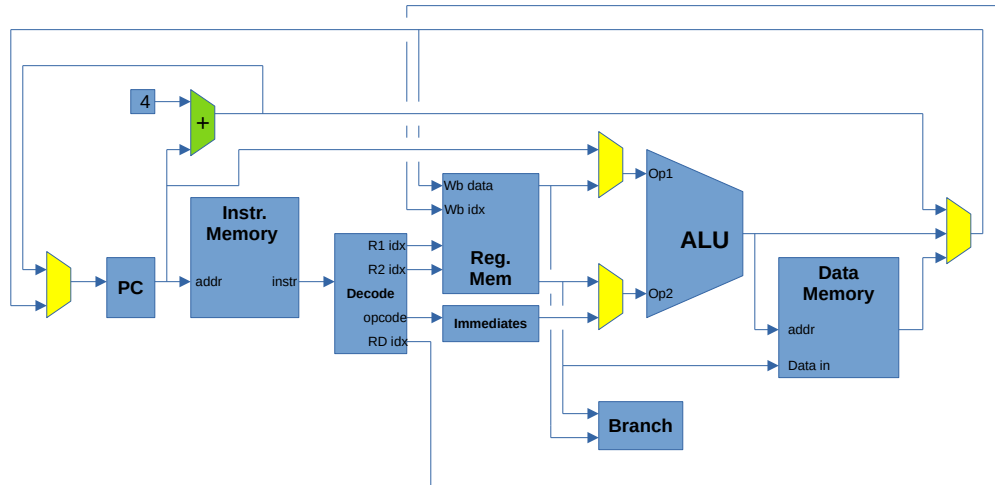


Figure 1.1: Basic CPU Design.

Most of the components shown here have already been explained earlier, albeit with some modifications and additions. It should be noted that the components are arranged from left to right in the order of the stages; with the wires of the WB stage looping back to where the outputs should be stored. The notable changes are as follows:

1. The decoder is actually two components: The Decoder and the Immediates Decoder. The Decoder focuses on getting the indices of the two input registers and the destination register, and then a set of values called the opcode, funct3, and funct7. These are 8-bit, 3-bit, and 7-bit numbers respectively that, together, describe which exact instruction it is. The Immediates Decoder gets the opcode in order to get the immediates out of the instruction; depending on the exact instruction, the immediate is found in different bits; RR instructions have no immediates at all.
2. There is a Branch Comparator component alongside the ALU. When executing a branch instruction, this component calculates whether the jump should happen or not while the ALU is busy computing the final destination of the jump.
3. Lastly, there are multiplexers, drawn as yellow trapezoids, throughout. These simply take multiple inputs and, based on the state of the CPU, decide which one should pass through. For example, the two multiplexers right before the ALU need to decide to let the values from registers r1 and r2 pass on an RR

instruction, but let the current PC and the Immediate pass in during a Control Transfer instruction.

Note that there are additional data pathways, going to the multiplexers specifically, that they use to determine which value should be allowed through. The leftmost "Jump" multiplexer checks the instruction's opcode and the Branch Comparator to see if it should jump, for example. These wires do exist in the real CPU, but are not depicted in the diagram for the sake of clarity.

This is enough to understand the CPU simulator, for the most part. Rather than simply taking in a machine-code program and execute it as efficiently as possible, which would be relatively trivial to do in a modern high-level programming language, the simulator represents each of these components with individual variables and updates them the same way a real CPU would, so that the flow of data through the CPU circuits can be displayed accurately.

Chapter 2

CPU Design

Before looking at the actual CPU simulator program, it is worth going over one particular feature that this simulator seeks to implement.

2.1 Pipelining

The previously-described CPU design has one simple issue: it is painfully slow. This design is capable of executing one instruction per cycle: one instruction per full use of each of the CPU's parts. This may seem good, but at any given moment, most of the CPU is actually going unused: the decoder is waiting on the instruction memory to fetch the instruction it needs to decode, which is waited upon by the register memory that needs the register indices to fetch their data, which is waited upon by the ALU that wants its inputs, and so on. When, say, the ALU is computing, most of the other components of the CPU are all entirely dormant. They could, theoretically, be doing something useful at the same time. However, because our design handles one instruction at a time, the components before the ALU have already done their part and the components after the ALU are waiting on it to finish and present its output. If one could find a way to utilize all components of the CPU at the same time, they could make a CPU with the same hardware perform many times faster.

This is exactly what Pipelining does. Pipelining is not so much a feature of the simulator as much as it is a feature of the CPU itself that this program simulates, and the design of the CPU itself must be significantly altered in order for it to be successfully implemented. However, this feature is quite powerful, and used near-universally in modern CPU designs ¹, and so the simulator is made with it in mind. One could describe the pipelined design as this simulator's main feature.

¹Denning (1993)




Figure 2.1: Program execution over time.

The solution is simple in theory, though difficult in execution: have the CPU process different parts of multiple instructions at once. Five of them, to be exact. This is where the five stages come into play: after each stage performs its function, it accepts the next instruction right away rather than waiting for the instruction to finish passing through the other stages. If the WB stage is executing instruction number 9, the MEM stage is executing instruction 8, and the EX is executing instruction 7, so on.




Figure 2.2: Program execution over time, with pipelining.

Remember that the choice to subdivide this process into five stages, specifically, is entirely arbitrary. It's easy to see how two of these stages could be combined to have less-than-5 pipeline stages, or how the Instruction Decoding and Register Fetching tasks could be separated into their own stages, but you can go much further than that: While 4 stages was the norm just three decades ago ², the now-also-outdated *Pentium 4 Prescott* CPU had a whopping 31 stages ³! And, by contrast, many newer CPUs like AMD's *Jaguar* and Intel's *Silvermont* have found a comfortable middle ground

²Denning (1993)

³Shimpi & Wilson (2004)

somewhere in the 10-15 pipeline stage range ^{4 5}. Theoretically, you'd want to split your CPU into as many pipelined stages as you possibly can; the more instructions you're running at once, the more efficiently the CPU is being utilized. There are tradeoffs and reasons for why this is not the case, which will be described later in this section.

Still, this is much, much faster; in optimal conditions, the CPU can execute nearly 5 cycles in the time it would've taken it to do 1 cycle before! Sure, each of those cycles only performs one-fifth of five instructions, but the end result over a long program would still be a five-fold increase in speed!

Of course, conditions are not always optimal. A certain addition will be needed to this diagram for pipelining to be possible at all, a few smaller ones will be needed to handle new issues that arise out of pipelining specifically, and some of those new issues can only be mitigated, not avoided fully.

The first such modification, which makes pipelining possible on a basic level, is to handle coordination. Theoretically, if every subcomponent of a CPU was perfectly coordinated and each of the five stages each always took the exact same amount of time to complete, one could simply set the CPU going, with a slight alteration to the PC so that it increments its program count 5 times per cycle instead of once. However, this is not the case: in real hardware, one cannot rely on parts being perfectly coordinated; and some stages will take more or less time than others.

The Memory stage, notably, can require time to perform that is magnitudes longer than the other stages when it is required to go off the CPU and into a separate RAM stick, which is needed due to the sheer extra amount of information that the Data Memory is expected to hold. Data Memory structures even tend to have specialized on-CPU caches to minimize the need for this; it was originally intended for the CPU simulator to represent these fully as well, but this was proven to be beyond the scope of the project. Instead, the Data Memory is abstracted as a "black-box" subcomponent, which stores an array of data accessible by address through unknown means. Indeed, this entire timing aspect will be largely abstracted away in the simulator, which operates entirely on the time unit of a "step". A step is an arbitrary unit of time, defined as the amount of time needed for all of the instructions currently in the CPU advance by one stage. However, it should be noted that the modification used to solve this timing issue in real hardware is essential to the CPU design, and is also used in some of the solutions for smaller problems that cannot be abstracted away. Therefore, it will be included in the simulator regardless.

2.1.1 The Backbone of Pipelining: Latches

So, what is this modification? To keep all five stages in sync, no stage of the CPU should receive the inputs for the next instruction before it is done performing the computations done on the current instruction. The ideal way to do this is to "hold" the outputs of each stage until the slowest stage completes its computations, and then

⁴Kanter (2014)

⁵Kanter (2013)

release them all at once to begin the next step. This "holding" of outputs is done through the addition of a set of objects called latches to the CPU:

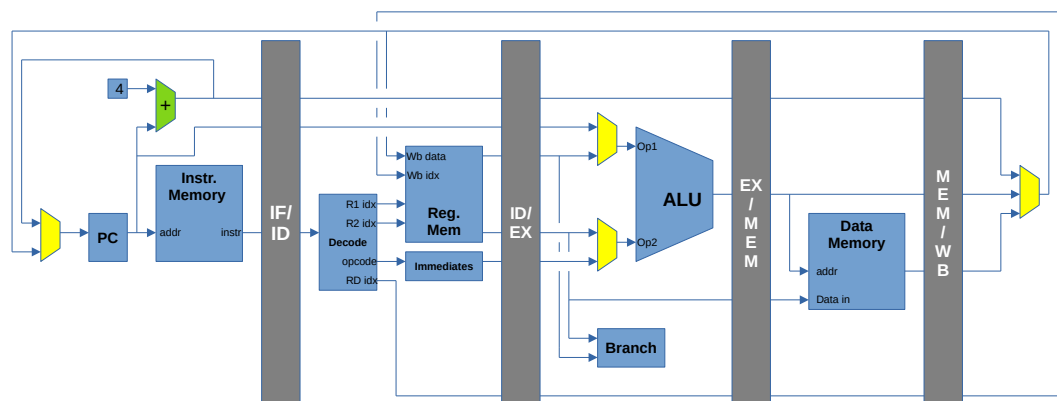


Figure 2.3: CPU with Pipelining Latches.

There is one latch in-between each CPU stage, with all of the outputs passing through said latch. Each latch is named after the two stages they separate, with the leftmost being the "IF-ID Latch" for example. No latch is needed after WB, since that stage simply consists of delivering the instruction result to where it needs to go. The latches are their own kind of register; in fact, "Pipeline Register" is another name for them. The latch holds one value for every wire that would pass through it, and constantly outputs said value through to the output end. Then, at the start of every cycle, the latch checks input end and updates its held value to match it, "opening" the latch. If the ALU's output wire read "13" before the EX-MEM latch was closed, and "4" after, the latch will continue to output "13" to the wire leading from the latch to the Data Memory until the latch is opened again. This achieves the 'holding' effect that was desired, and makes pipelining possible at all!

This is one of the reasons why having more pipeline stages incurs diminishing returns, along other tradeoffs. Simply put, latches take a small amount of time to open, during which no computation can occur. The more pipeline stages your CPU has, the more times this needs to happen during one cycle, and the longer the cycle gets; therefore, adding more pipeline stages inherently has diminishing returns for the CPU's efficiency. There are entire studies about measuring the efficiency of CPUs, and how this is impacted by design decisions such as its latch placement.

However, even with pipelining and latches in place, even with the efficiency trade-offs in mind, there are still other issues that will arise from this feature. Ultimately, whoever wrote the code that this CPU is running did so under the fundamental assumption that the CPU executes one instruction of code after the other, sequentially, waiting to finish one instruction before beginning the next. This is not the case in our faster, pipelined CPU; and while that's fine in most situations, it sometimes isn't. There are certain scenarios, called "hazards", that can occur in a machine code program where, due to this false assumption, our pipelined CPU as-is will produce an output that is wildly different from what the program correctly should. In order for

the pipelined CPU model to be valid, it must have some way to avoid or correct any hazards that come up! Hazards can be broadly divided into these groups:

1. Data Hazards: Hazards that involve the reading and writing of data to registers.
2. Control Hazards: Hazards that involve jump instructions.

Let's look at these hazards in more detail, and at how the pipeline CPU design can be altered to handle them.

2.1.2 Data Hazards

Data hazards are best introduced with a simple example, shown in assembly code. Assembly Code is only one step removed from Machine Code: It's the same line-by-line instructions that deal directly with the CPU's registers and the Data Memory, for the most part, but presented in a way that is somewhat human readable instead of as rows of bits.

For this example, all you need to know is the RISC-V instruction ADDI. The instruction "addi b, a, i", where *a* and *b* are registers and *i* is an immediate value, adds the contents of the *a* register to the immediate *i* and stores the result in register *b*. It is not unlike the line of code "b = a + 10" in a python program.

Imagine the following program snippet, where \$r0, \$r1, and \$r2 are registers and all registers start with the value 0.

```
addi $r1, $r0, 3
addi $r2, $r1, 4
```

In a non-pipelined CPU, this program runs fine, as shown below. The steps relevant to this example are described by cycle; the IF and MEM stages do not matter to the example and will thus be skipped over.

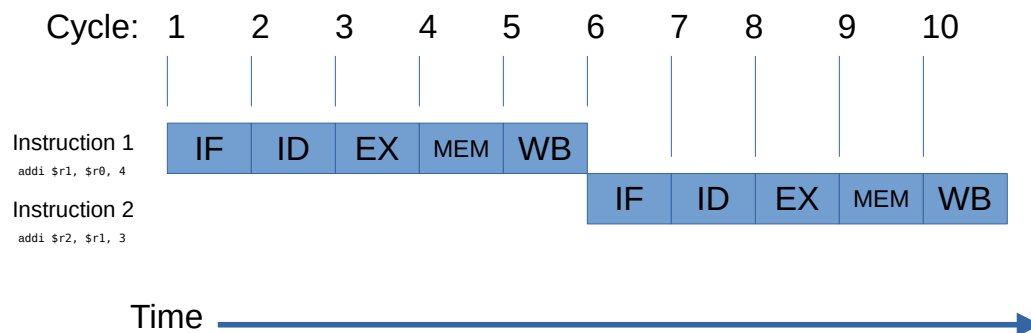


Figure 2.4: Program execution without pipelining.

Cycle 2: Fetch the value of input register \$r0, 0, from Register Memory.
 Cycle 3: ALU executes $0 + 3$, its output is 3.

Cycle 5: The output, 3, is written back to the $\$r1$ register in the Register Memory.
 Cycle 7: Fetch the value of input register $\$r1$, 3, from Register Memory.
 Cycle 8: ALU executes $3 + 4$, its output is 7.
 Cycle 10: The output, 7, is written back to the $\$r2$ register.

The end result is a state where $\$r1$ holds the value of 3 and $\$r2$ holds the value of 7. This is the correct output, and what any such person would expect this code snippet to perform with the described starting conditions. However, look at what happens if we run the same code snippet through our pipelined design:

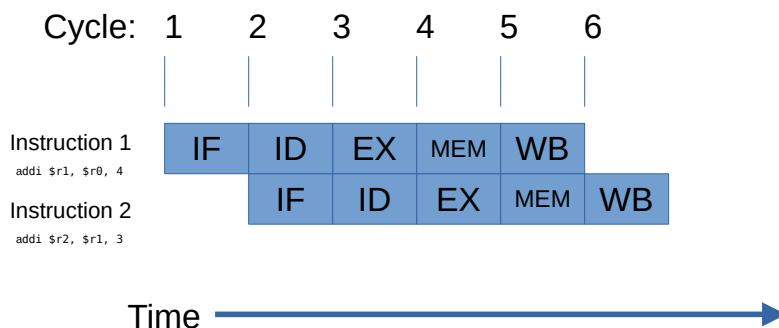


Figure 2.5: Program execution with pipelining.

Cycle 2:

Instr 1 ID: Fetch value of input register $\$r0$, 0, from R Mem.

Cycle 3:

Instr 1 EX: ALU executes $0 + 3$, output is 3

Instr 2 ID: Fetch value of input register $\$r1$, 0, from R Mem.

Cycle 4:

Instr 2 EX: ALU executes $0 + 4$, output is 4

Cycle 5:

Instr 1 WB: The output, 3, is written back to the $\$r1$ register.

Cycle 6

Instr 2 WB: The output, 4, is written back to the $\$r2$ register.

This time, the end value of $\$r2$ is erroneously 4 instead of 7! The pipelined CPU gives an incorrect output, and the reason boils down to this: **the output of Instruction 1 is used as input in Instruction 2, since $\$r1$ is both the former's destination register and the latter's input register.** But, Instruction 1's WB stage happens **after** Instruction 2's ID stage. In most cases, this would be fine; but not when the second instruction is trying to read from the same register that the first instruction modifies. It will happen even if these two instructions aren't directly sequential; only after the two are far enough temporally that the former's WB stage occurs at the same time as the latter's ID stage does it stop being an issue.

At the very least, having a program return a consistently incorrect output is

unnacceptable, and we need to fix that. But, how?

2.1.3 Stalling

The first solution that might come to mind is: "Why don't we pause the CPU when a Data Hazard is incoming, so that it has time to catch up?". And, this is shockingly a decent suggestion! The idea is simple: As it is in the example, the update to `$r1`, occurring at the start of Instruction 1's WB stage, was needed halfway through Instruction 2's ID stage (Register accessing happens during the latter half of the ID stage, after the Decoder has gone). The red arrow in the diagram represents this desired update that needs to happen for correctness.

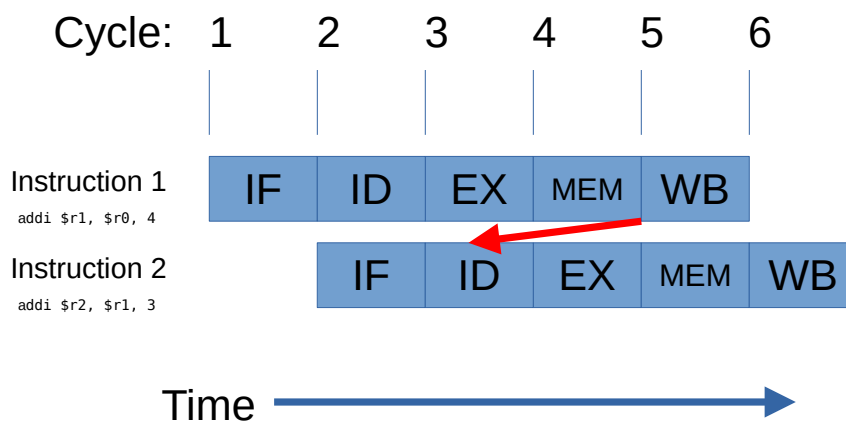


Figure 2.6: Program execution with no Stalling. Data Hazard occurs.

So, we could just delay the execution of Instruction 2 by three cycles...

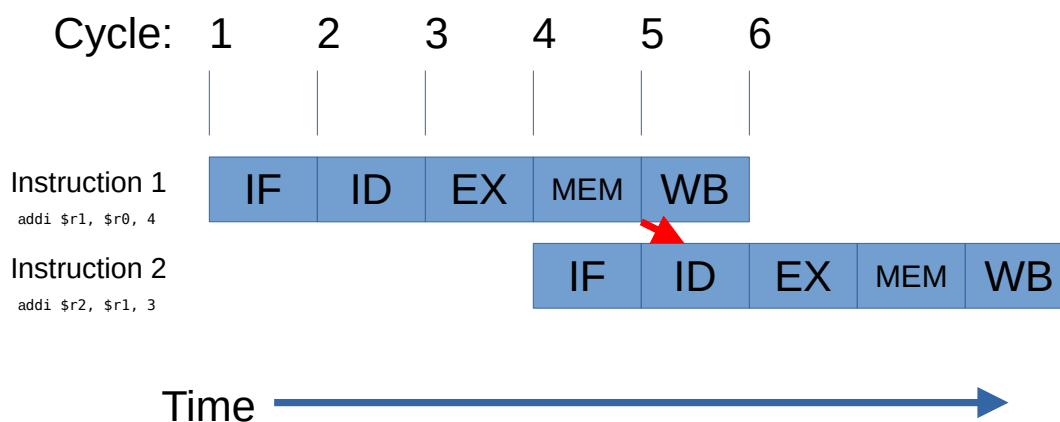


Figure 2.7: Program execution with Stalling. No Data Hazard.

And now the program will execute correctly! This process, called "Stalling", can be applied dynamically, too; if a Data Hazard happens by just one stage of difference,

the offending instruction only needs to be stalled one cycle. This method inherently costs us some of the efficiency that pipelining offers either way, though.

But, how is this Stalling implemented? As it is now, our pipelined CPU functions, but it will simply barrel forward without control: At the start of every CPU cycle, all the latches fire and transfer their information forward to the next stage. Luckily, we already have the perfect tool to control this flow, and it's those latches themselves!

Let's add a new functionality to our latches. As they are now, the latches are in "Transfer" mode:

Transfer: At the start of cycle, take the incoming currents of the previous stage and update the outgoing currents to match.

Now, let's add two other possible modes:

Stall: At the start of cycle, change nothing; keep the previous outgoing currents as they were.

Bubble: Set the output values to all-zeroes, regardless of what the input values were.

Lastly, we add a "Stall Control" subcomponent to the CPU. All it does is determine whether and how it needs to Stall for the given cycle, and change the latches between these three states to perform the stall. Note: There is an additional fifth "Latch" in this setup, where the Stall Control hooks up to somewhere in the Program Counter circuitry. This simply allows the Stall Control to pause the Program Count's natural per-cycle increment; so long as the instructions are stalling, the program count needs to stall as well or the CPU will start skipping instructions entirely.

But, why are all these modes needed at all? Well, we cannot simply stall out the entire CPU; that is the same as if it wasn't running at all. What we need to do is allow for some instructions, those before the hazard-bearing instruction, to finish running while the rest of the program stops and waits for the desired number of stall cycles.

The Stall Control has all latches set to "Transfer" during normal operation. However, once an instruction with a Data Hazard reaches the ID stage, it alters the latch modes as follows.

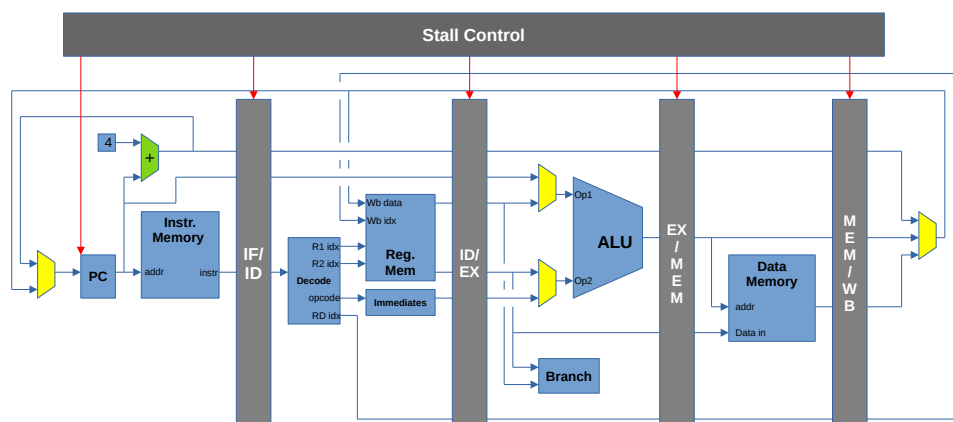


Figure 2.8: CPU Design with stall control.

At this point, the EX, MEM, and WB stages hold the instructions that you do not wish to stall; those are Transferred so they can finish. The IF and ID stages are Stalled (using the PC and IF-ID Latch respectively), holding their current instructions. Lastly, the EX stage gets "Bubbled"; set to an output matching a "NOP" instruction that simply does absolutely nothing. This is quite important: if the EX stage was simply stalled as well, you'd actually start getting duplicates of the instruction immediately preceding the hazard-holding instruction inserted into the pipeline. Thanks to the bubble, harmless 'nop' instructions that alter nothing get sent out instead to fill that gap.

Once the instruction that holds the needed alteration reaches the WB stage, the Stall Control simply flips all latches to Transfer mode, and operation continues as normal. This does successfully resolve the data hazard, and fixes these correctness issues! However, it is still painfully slow, and gives up much of the advantage of pipelining anytime that a register is written to and then read from in quick succession; as it turns, this is rather common. We've already accepted that conditions will not often be optimal for pipelining, but this is a serious blow. Can we do better than that?

2.1.4 Forwarding

The answer is that yes, we can! Let's go back to our data hazard example, without the stalling:

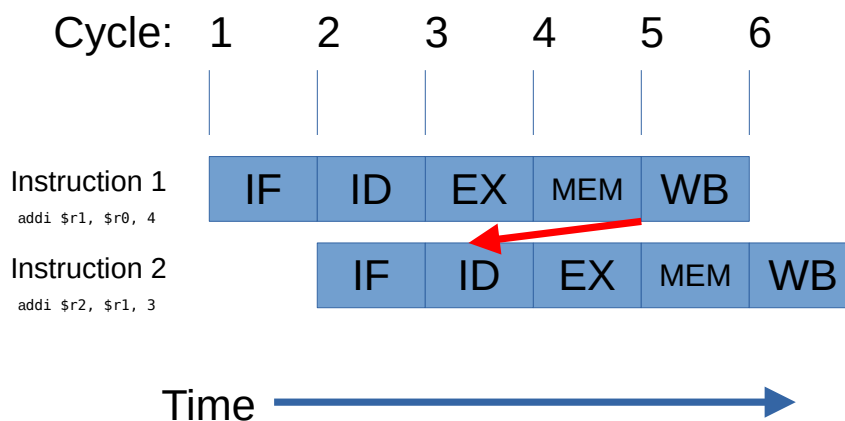


Figure 2.9: Program execution without Stalling

Obviously, the issue is clear here: The data needed in the ID isn't ready until Instruction 1's WB stage, and is needed in Instruction 2's ID stage. Except, this isn't really true. Technically, that data has already been computed by the end of Instruction 1's EX stage, found in the output of the ALU. And, it is Instruction 2's EX stage that needs the data for the addition calculation, not the ID; as long as the ALU gets the correct integer input of 3 instead of 0, it doesn't care where it came from. So, technically, the desired data to resolve this hazard *is* available by the time

it's needed; it just doesn't have the pathway to get there in time. So... what if we added that pathway ourselves?

That is exactly what forwarding is! Extra data pathways are added into the ALU input; controlled by a multiplexor that reads register indices to decide if forwarding is needed. The pathways come from the MEM-stage output if the offending instructions are adjacent, and from the WB-stage output if they are one instruction apart; any more than that, and there is no Data Hazard at all. These are called "EX-EX Forwarding" and "MEM-EX Forwarding" respectively. There is a third type of Data Hazard, which occurs when a value is loaded from the Data Memory to a register and then stored to somewhere else in the Data Memory immediately after; the store will still have the stale value from that register, from before it was read to. An extra pathway and multiplexer, from the WB stage to the Data Memory, handles that; this is called "MEM-MEM Forwarding". With that, all of the Data Hazards are resolved, this time without performance loss!

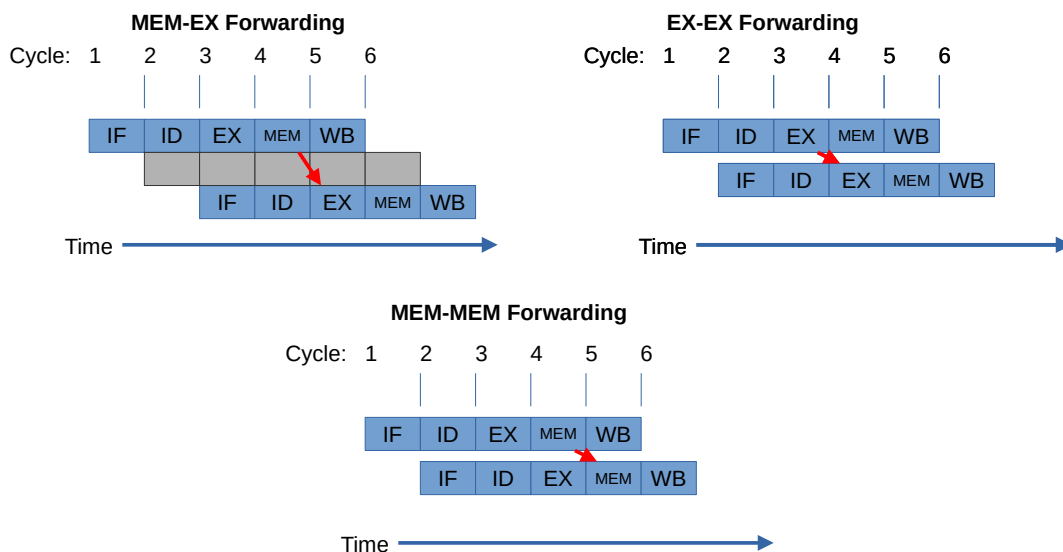


Figure 2.10: The three major types of forwarding.

Doesn't this make stalling entirely obsolete? Well, not exactly; there's still one more unhandled data hazard case. What happens if a load instruction is immediately followed by an arithmetic instruction, using the register that was loaded to? (Or by a store instruction, using the register to calculate an address rather than as the to-store data). The data is retrieved at the end of the Load's MEM stage, yet is needed at the beginning of the EX stage so it can be used in the arithmetic operation. In this case, the data truly doesn't come soon enough!

The ideal solution is a hybrid of the two techniques: If you stall the second instruction by just one cycle, the data can now arrive just in time. It can even use the same "MEM-EX Forwarding" from earlier without any new pathways needed!

2.1.5 Control Hazards

There is a second reason why stalling, or at least the Stalling Control we implemented, is still useful: Control Hazards.

Control hazards occur anytime any jump instruction is called. The problem is simple: anytime that a jump instruction occurs, the new address to jump to is already needed by start of the next instruction's IF stage, but it is calculated in the jump's EX stage, where the ALU adds the jump value to the existing Program Count.

This hazard is going to lose us cycles no matter what. To mitigate that, the first step is to add another forwarding pathway from the ALU output directly to the PC in the IF Stage. This almost resolves the issue, with the jump happening at a 2-cycle delay, but there's one more problem: in between the jump instruction entering the IF stage and the jump address being calculated in its EX stage, two instructions directly after the Jump get loaded in; instructions that shouldn't run at all, since the program supposedly jumps to another location. After those two, though, the output of the EX stage gets forwarded back to the PC in the IF stage, so the solution is simple: On the cycle that the jump instruction is in the EX stage, we tell the Stall Control to set the latches of the IF and ID stages (the stages holding the two erroneous instructions), to go on "Bubble Mode". This harmlessly flushes out those two instructions before they have a chance to perform any alterations, and the instruction following them will be from the correct jumped-to address! The end result is akin to stalling the instruction following a jump for 2 cycles, in order to avoid the Control Hazard.

In fact, you could achieve the same result by stalling the instruction 2 cycles, using the method described in the earlier section on stalling. This would work: there, the two erroneous instructions never enter the pipeline at all. It would be simple if all hazards were handled through the same method, so why not just do that?

It has to do with *Branching* control instructions. A Branch is similar to a jump, but it also makes some sort of logical comparison between two inputted register values. If the comparison resolves to true, the Branch jumps to the destination; if it's false, the jump is not taken and the program continues to the next sequential instruction. Branches are the main option assembly programs have for introducing, well, branching pathways to the control flow of their code, and what compilers most heavily use when compiling if-statements and loops from high-level programming languages.

Well, as it turns out, handling jumps with the flushing method is actually a small optimization that potentially saves time for *branching* control instructions. Branching control instructions are almost identical to jump instructions: there's just an extra subcomponent in the EX stage that compares its values to check if the Branch is to be taken, which decides if the Program Count is changed or left alone. However since we cannot know whether the branch will be taken or not *until* the EX stage, this causes a difference between the flushing and stalling methods: In the stalling method, the instruction after the Branch has to get stalled regardless of whether the Branch gets taken or not; it has to start stalling right away. In the flushing method, the CPU flushes and jumps as usual if the Branch gets taken; however, if it does not get taken, the Stall Control simply... doesn't do anything. Those two "erroneous instructions" are actually now the correct path forward for the program to take, and

it just lets them happen. So, flushing the Branches instead of stalling them makes no difference if the branch is taken, but saves cycles if it isn't! This method is also called the "Fetch and Cancel When Taken" method of handling control hazards, since the "cancel" or flush only happens if the Branch is taken.

2.2 Final CPU Design

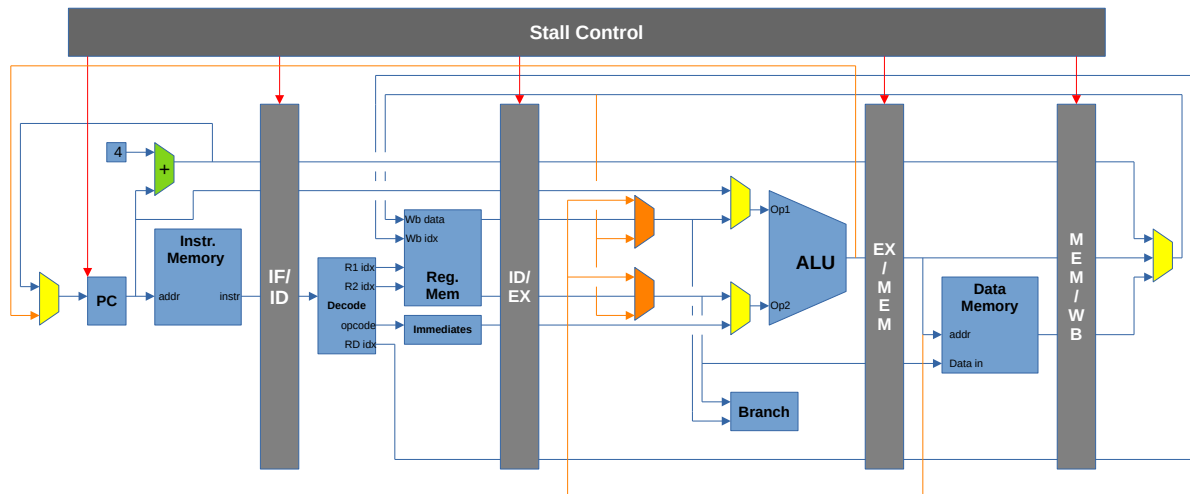


Figure 2.11: Final CPU Design.

The final design of the CPU we will simulate is shown above. Instruction, Decoder, Registers, Arithmetic, Data Memory, and Writeback; built with the intent that the simulator will display its own internal state to a similar level of granularity as the diagram above.

The flow of instructions through the wires is separated into 5 arbitrary stages in order to **Pipeline** them for major performance boosts, with the separation enforced by a set of Latches.

Then, in order to deal with the correctness hazards that this pipelining incurs, additional **Forwarding** pathways and a **Stall control** in the latches are implemented to sidestep those hazards, while also minimizing the performance loss incurred. The new forwarding pathways are highlighted in orange in the diagram above. Note that two additional multiplexors, one for R1 and another for R2, are needed in order to determine whether a jump is needed there at all.

Chapter 3

Simulator Design

The actual simulator is a program written in the Rust programming language, and aims to faithfully replicate the operation of the CPU at the same level of abstraction used in the diagrams above! As a result, it can be compiled and run on most modern personal computers. Its final iteration fully supports the Risc-V 32-bit Base Integer Instruction Set, with the exception of system-call instructions that are meaningless without the context of an OS outside the program, and the failsafe for a particular Data Hazard that turned out to be rather difficult to implement.

The goal of the simulator itself is simple: to take in a machine-code RISC-V program as input, set up a virtual construct in Rust that replicates the low-level, semi-abstracted model of a CPU that has been discussed in the background chapter, and then run said program through it in a way where the simulated CPU's internal state is always transparent. Other goals, such as a User Interface, are additional features and not part of the core design. This design was eventually accomplished, though the project's initial design needed to be discarded and a new design made anew.

3.0.1 First Draft

The first draft was quite neat in theory, though rather idealistic. The core idea is to rely on the way that the diagrams above divide the CPU into separate "subcomponents", such as the Program Counter or the ALU, which only pass along information through predefined wires; this is something the final design will keep.

The core idea of the first draft is that each of these subcomponents should be represented by its own unique "struct" in Rust; despite the name, these are much closer to the Classes and Objects of Object-Oriented Programming Languages like Python and C++, made to hold a specific set of variables as well as functions tailored to those variables. Every "subcomponent" on the graph would be represented with its own unique struct. It would have variables for both every input and output connection it had with other subcomponents; and if it was a stateful subcomponent that stores data, like the Program Counter or any of the three Memories (Instruction, Register, Data), it would also have variables for recording said state. References to the structs that they need to fetch input data from, as well as functions to fetch said data and

to perform the subcomponent's intended purpose complete the struct. Then, all of this would be hooked up to some sort of master "setup" and "step" functions; the former would initialize all these structs and insert the program into the Instruction Memory, while the latter would activate these structs one by one to perform a single full CPU cycle, over and over.

The advantages of such a design are apparent: The different structs are entirely separate from one another, making them easy to test on an individual level and to see the state of any one given subcomponent at a time. It also meant that the layout of the code resembled the design of the actual CPU somewhat intuitively.

However, this design also has two glaring flaws: The first is more manageable, and it is simply that the design's method of passing data between structs, that being of giving each struct a persistent reference to the structures that send information to it, is entirely incompatible with Rust's memory-safety failsafes. Well, the idea of doing that, in itself, is not, but the problem comes from the fact that, due to the CPU's cyclical nature, these links of one struct borrowing a reference to the previous struct eventually form a loop. The details are not relevant to this project, but using such persistent references in a loop with no clear top-level "owner" goes against the language's philosophy to the point of not being allowed at all. This issue could've been resolved on its own if it was the only problem with the design; for example, by simply having the step function handle all the wiring itself and using the subcomponent structs for their internal computations only.

That said, the second issue is that the design was unbearably repetitive, to the point of being confusing to look at. The distinct input-output variables on each struct meant that every piece of information in the CPU essentially took up two variables instead of one, and the process of computing a piece of data and sending it to the next subcomponent in the stage became a chore with many unnecessary middle-man steps inbetween.

3.0.2 Final design

The second and final iteration of the simulator took a slightly different approach, one that relies more closely on the nature of this CPU as being ultimately made of logic gates and wires, and employs style of design often used to simulate simpler logic gate machines.

The idea now is to divide the CPU design into two categories, and therefore two structs: Its State and its Logic. The State consists of all the naturally stateful elements of the CPU as described before: The Program Counter, the three Memories, and also all of the Latches, which essentially store the supposed current state of the wires they are supposed to clamp down on. The Logic stores everything else that has no state; or, rather, it stores all of the *wires*, since that is all else that matters. A subcomponent like the ALU has no internal state; it reads data of the two operands and the opcode, and calculates its output solely off of that; even to be accurately represented, this representation can be done simply as a piece of code that performs the ALU's function. Even the wires are only stored in persistent variables for the sake of visually displaying the CPU's state later. So, the entirety of the CPU is

separated into a State struct and a Logic struct, and each struct has an "Update" function written to it that takes the other struct as a mutable reference and updates the state of all its own variables at once: For the State Struct, this means handling memory changes like updating the PC, enacting the write of data into the Writeback register, and Storing data into the Data Memory when Store instructions pass by. For the Logic struct, it means reading the state of latches and registers and determining what signals should be traveling through the wire as a result of their changes.

Note that some stateful components still have stateless outputs. Here, the component stores a state, but the output wires are entirely dependent on the stateless input; therefore, the output is stateless as well. Wires are only stored in the State struct if they have an output that's solely dependent on the component's state, such as the PC which always displays its internal value or all of the latches' output.

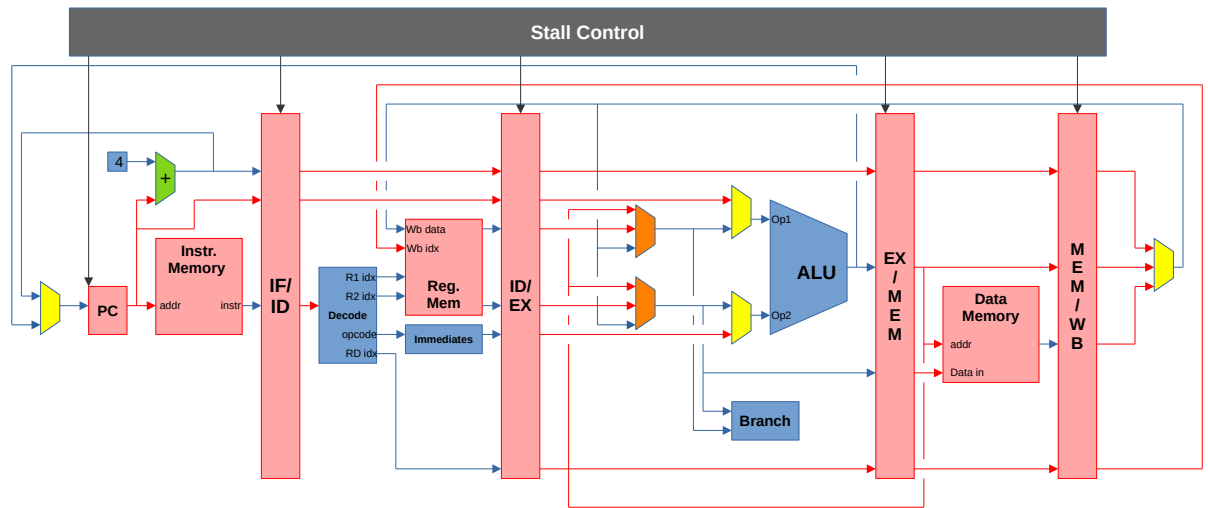


Figure 3.1: Color-coded CPU design. Red is for data stored in the State struct, Blue is for data in the Logic struct. Multiplexors are definitionally stateless, and are all in the Logic struct.

This setup is so useful because it makes the actual update of a CPU cycle incredibly simple: Just run the update function for the State struct, followed by the update function for the Logic Struct. This is not entirely true, as the implementation of stalling will complicate things a little, But for the most part, this makes up the base of the simulated CPU and the process for executing one of its cycles:

List of Listings 3.1: Barebones layout of *components.rs*

```
struct State {
    pc : u32,
    instr_mem : Vec<u32>
    // ... more State variables
}
```

```

struct Logic {
    added_pc : u32,
    instruction_out : u32,
    //... more Logic variables
}

impl State {
    fn update(&mut self, &mut logic: Logic){
        self.pc = logic.added_pc;
        //... more State updates
    }
}

impl Logic {
    fn update(&mut self, &mut state: State){
        self.added_pc = state.pc + 4;
        self.instruction_out = state.instr_mem[(state.pc / 4)]
        //... more Logic updates
    }
}

fn step(state: &mut State, logic: &mut Logic) {
    state.update(logic);
    logic.update(state);
}

```

This code is entirely functional as a rough base for the simulator's design; all the main body of the program needs to do outside of this function is to initialize the two structs, get the user-fed instruction into the instruction memory, and then run the *update(state, logic)* function repeatedly until the program has run its course!

Other features can and will be implemented as well, of course; but, the design shown above is a sweeping generalization. It's worth taking a closer look at what the final, detailed program looks like, now that its basic structure is understood.

3.0.3 Structs and Sub-Structs

The final design could indeed have just unceremoniously dumped all of its State and Logic variables into their respective structs; but this would've made a rather huge mess, resulted in individual wires (especially different wires that carry the same piece of data at different stages) being hard to keep track of, and would've been rather confusing and inelegant in general; this is one of the issues the first draft design sought to resolve. Therefore, the final design for the structs instead looks like this:

List of Listings 3.2: Snippet of *components.rs*, holds main 2 structs

```
#[derive(Clone)]
```

```

pub struct Registers {
    pub ifid: IFIDLatch,
    pub idex: IDEXLatch,
    pub exmem: EXMEMLatch,
    pub memwb: MEMWBLatch,

    pub pc: u32,

    pub instr_mem: Vec<u32>,
    pub reg_mem: Vec<u32>,
    pub data_mem: HashMap<u32, u32>,
}

//this struct holds all the wiring of each stage

#[derive(Clone, Default)]
pub struct Logic {
    pub fetch: IFLogic,
    pub decode: IDLogic,
    pub execute: EXLogic,
    pub memory: MEMLogic,
    pub writeback: WBLogic,

    pub pc_stall: bool,
}

```

Rather than having a soup of variables, all of the states tied to a specific latch and all the wires tied to a specific pipeline stage are sequestered off in their own sub-struct; when used like this, these sub-structs are just fancy ways of holding variables and don't get in the way like they did in the previous design. It also allows for said sub-structs to be stored in their own file, leading to a much cleaner folder structure for the source code. In a way, these sub-structs are effectively just organizational folders for the variables within the program.

List of Listings 3.3: *decode.rs*, the file that holds ID-stage structs

```

use crate::components::*;

//The IF-ID Latch
#[derive(Clone, Copy, Default)]
pub struct IFIDLatch {
    pub base_pc: u32,
    pub added_pc: u32,
    pub instruction: u32,

    pub id_stall: u8, //stall, bubble, or neither?
}

```

```

impl IFIDLatch {
  pub fn bubble(&mut self) {
    self.base_pc = 0;
    self.added_pc = 0;
    self.instruction = 0;
  }
}

//holds all of the wiring of the ID stage
#[derive(Clone, Default)]
pub struct IDLogic {
  pub decode_r1: u8,
  pub decode_r2: u8,
  pub decode_opcode: u8,
  pub decode_rd: u8,
  pub decode_funct3: u8,
  pub decode_funct7: u8,

  pub regmem_r1: u32,
  pub regmem_r2: u32,

  pub immediates: u32,
}

```

This is an example of one such file; in general, everything is stored as simply as possible, with the data types one would expect: Most wires carry a 32-bit unsigned integer; since this is a 32-bit system, that is the format that the CPU refers to handle its data in while wires for the smaller opcode, funct-3, and funct-7 all use unsigned 8-bit integers. Technically, the latter should use u7, u3, and u7 respectively, but establishing custom-length integer data types just for them was deemed to not be a very high priority.

The instruction memory and register memory are simply represented as u32 Vectors, with the former having a length equal to the number of instructions in the program and the other having 32 integers exactly; this makes accessing registers incredibly intuitive, as the *n*th register can be found in *state.reg_mem[n]*. The data memory is orders of magnitude larger than the first two though, and is going to consist mostly of empty addresses, so using a vector would have been incredibly wasteful; instead, it uses a hashmap of 32-bit integers with numbered keys, matching the size of a "word" that the *Store Word* and *Load Word* instructions use. This means that the smaller load-store instructions need to slice off part of the accessed word, but this was fairly manageable.

It is worth noting that there are some wires which aren't often shown in CPU design diagrams, but are nonetheless present in a real CPU, and therefore in our simulation too. For example, the Forwarding Multiplexors are only shown as having

the base, EX-EX, and MEM-EX register values going into them, but in reality they also read a few other values, such as the opcode and register indices of the current and next two instructions, to decide which forwarding, if any, should occur.

Ultimately, this makes the storing and accessing of information about the CPU pretty intuitive, which will come in very handy for the User Interface.

3.0.4 Update Functions

Unfortunately, the two update functions aren't quite as simple. In a sense, they are: Each one goes through the struct going through all of its variables and updating them according to the internal logic of the CPU design. However, this updating needs to be done in just the right order in order to make things work properly. This is one case where building a virtual simulator is actually harder than making a real CPU: In a real CPU, all of this happens instantaneously and at the exact same time, literally within that given CPU clock cycle, and so it does not need to worry about this. But the simulated CPU runs one-step-at-a-time, and doing certain updates before others can cause massive errors. For example, the latches in the State struct's update function need to be updated in order from right to left, starting with the MEM-WB latch and ending with the IF-ID latch. Otherwise, if the IF-ID latch updates before the ID-EX latch, the information it was supposed to send directly over, like the base PC and added PC, that information vanishes for good. The problem is especially tricky with the Logic update function where, thanks to forwarding and writebacks, certain computations of a pipeline stage rely on information from a stage further to the right. Thankfully, part of the job of the latches here is to serve as a stopgap to separate the stages, so the issue is merely annoying rather than impossible.

A lot of the desired ordering was figured out, in part, by writing a large amount of test cases and catching any correctness errors, though that will be discussed in detail later.

Its ordering aside, the actual code contained within is mostly understandable, just there to perform the functions that those sub-components would in a real CPU. The decoder code takes an instruction and separates it into its relevant pieces using bit-masks and bit-shifts; the register memory writes and reads data off the 32 general-purpose registers, so on. The most crucial one is the ALU, since that's ultimately where the actual computation of most instructions takes place, and even that one is fairly simple: It takes in the two operands, looks at the opcode to figure out what sort of arithmetic operation it's supposed to perform and spits out the output to *logic.execute.alu_out*. To whomever is interested in more details, the update functions and my comments on them can be found in the *components.rs* file, right under the structs.

3.0.5 Stalling

There is one piece of the core simulator that doesn't fit neatly into the description above, in part because it operates above the other components. While the variables to indicate whether components are being stalled or not are stored within the State and

Logic structs, the code for the Stall Control happens directly at the step function, performed after the state update but before the logic. Since forwarding is implemented, stalling is only needed for Control Hazards and for Load-ALU hazards; the former is fully implemented, while the latter, unfortunately, could not be debugged in time.

```
pub fn step(state: &mut Registers, logic: &mut Logic) {
    //Checks if Stalling or Bubbling is needed
    //Check if a jump is performed! those always need a stall on the NEXT step
    let must_jump_stall_next = logic.fetch.jumped;

    //if there's a Load in EX stage and an ALU instruction in the ID stage,
    // using the same register...

    state.update(logic);

    //always do this inbetween!!!
    //that way the right instructions get stalled/bubbled,
    // but their logic is not allowed to propagate.
    if must_jump_stall_next {
        //If there is a jump, need to bubble IF and ID of now...
        // so the ID and EX of next step!
        logic.pc_stall = false;
        state.ifid.id_stall = 2; //bubble flag on
        state.ifid.bubble();

        state.idex.ex_stall = 2; //bubble flag on
        state.idex.bubble();

        state.exmem.mem_stall = 0;
        state.memwb.wb_stall = 0;
    } else {
        logic.pc_stall = false;
        state.ifid.id_stall = 0;
        state.idex.ex_stall = 0;
        state.exmem.mem_stall = 0;
        state.memwb.wb_stall = 0;
    }

    logic.update(state);
}
```

I suspect that this decision to place the stall decision there is part of what made the Load-ALU stalling so difficult to figure out, and is worth addressing in future revisions. But with that, the core simulator design is complete!

3.1 CPU-Simulator Features

Of course, such a simulator as it currently stands is pretty bare-bones for the sake of educational purposes. So, in a rough order of complexity-to-implement, the following features were added to the simulator: The ability to feed the simulator a machine-code program from a text file, the ability for the simulator to store and load the entire state of the CPU in previous steps, and a simple graphical user interface that displays and clarifies the internal state of the CPU at its current cycle. But, why these three?

The goal of these features are to make the simulator easier to use, and to aid its goal of helping users attain a more intuitive understanding of a CPU's inner workings. The last feature is by far the most critical for this: The GUI displays the same information as one of the CPU diagrams from earlier, but with the actual data of a live program running through it overlaid on top; and it can be analyzed step by step! The finalized GUI for this simulator is very rudimentary and much smaller in scope than what was originally intended, but it accomplishes its intended purpose well; it will be described in more detail later.

The purposes of the other two features are supplementary in nature, but no less useful: User-Inputed programs make it possible for the end user to provide their own code and run programs tailored around whichever instructions or CPU features they are interested in getting a better grasp of rather than one sample program, without needing to edit the source code and recompile at least. The latter is pure ease-of-use, and ends up implemented solely as a way to be able to "backstep" or "rewind" the simulator at any given step, but is incredibly useful in that alone: If a user steps forward and suddenly finds themselves confused by the CPU's new state, they can rewind back to the previous cycle, or more, to refresh their memory and see how it got there.

3.1.1 Implementation of User Input and Rewinding

These features were both pretty simple to implement, and are only described here for completeness. The User Input simply involves using Rust's built-in libraries for commandline arguments and file reading, taking in the filename and creating the Instructions vector out of the file's contents line-by-line. For the sake of simplicity, the program expects one machine-code instruction per line, formatted as a 32-bit binary integer. Originally, the project was going to accept RISC-V assembly code as input and have its own built-in assembler to translate them into the binary machine code that the simulator needs, but this has proven to be out of scope.

The backups/rewinding were only a little more involved, requiring just the creation of a "Snapshot" struct that contains one State struct and one Logic struct within it. This new struct then works as a full snapshot of the CPU's previous state: The simulator clones the current State and Logic structs at every cycle, and places the clones in a "backup" vector of Snapshots; When the user asks to go back a step, it simply pops the most recent Snapshot off the backup vector and replaces the State and Logic structs it is currently using with those within the Snapshot. Due to their simplicity, both of these features were very easy to implement!

3.1.2 Graphical User Interface

In the initial writeup, the plan for the simulator's GUI was incredibly ambitious (the entire writeup was, but that is besides the point): The plan was to learn an application development framework such as Qt as part of the project, develop a proper graphical interface that would run its own window, and then implement all sorts of additional nice features like separate tabs for the different memories, animations to show which wires are actually passing relevant information in any given cycle, and so on. Suffice to say, this idea was entirely out of scope and would have required spending a lot of time studying something entirely unrelated to Computer Systems, and I ended up scaling back to something much more reasonable to get working quickly: An ASCII-art text interface, that can be printed directly to the commandline terminal.

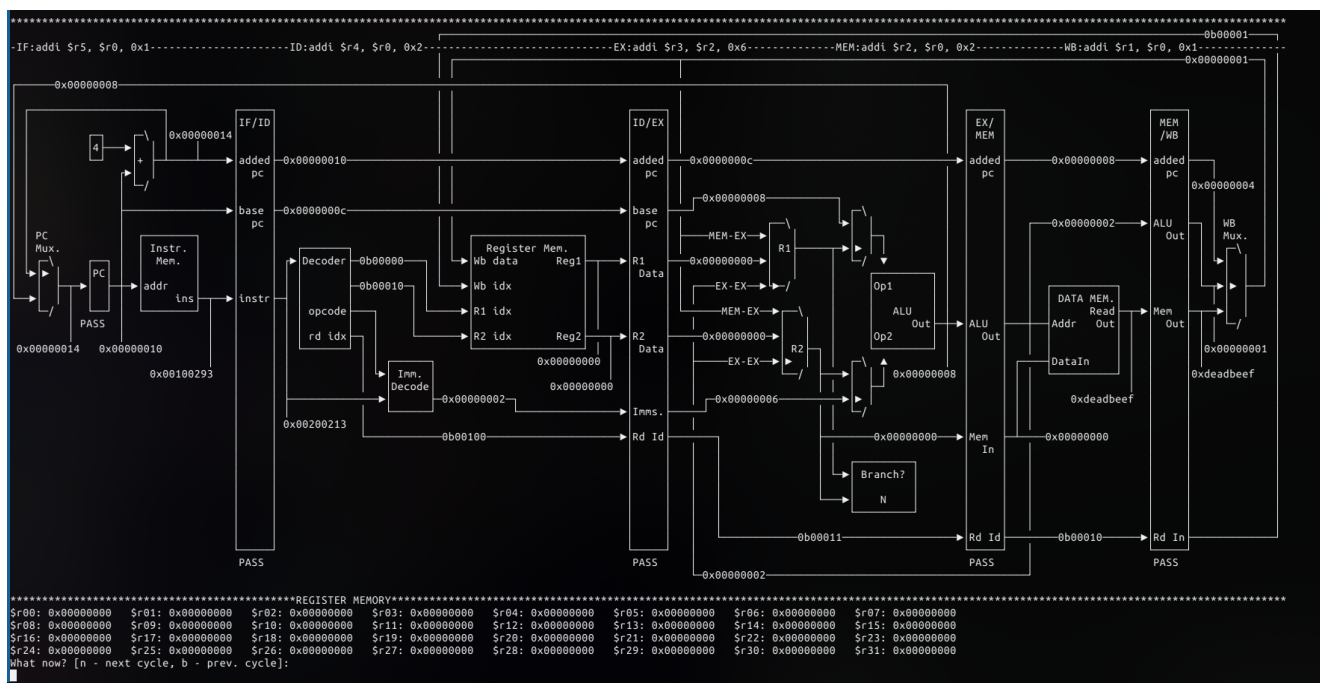


Figure 3.2: The Simulator's command-line-based GUI.

It's fairly dense and all but requires the user's terminal to be maximized, but it properly displays the CPU diagram and all of its relevant information! The top line displays the instruction that each pipeline stage is currently executing, in assembly; the body displays the entire CPU diagram, and the value that every wire currently holds; and the bottom shows the entire values of the register memory! The implementation, while involved and time-consuming, was mostly simple: the text-art for the diagram was drawn up using an online tool named ASCIIflow, and the *display_cpu* function responsible for it was given a reference to the State and Logic structs, from where it could effortlessly pull the needed values for the diagram display. This simpler method also provides its own advantage: While much less sleek, the commandline approach makes the program more likely to work across operating systems without any tinkering required than an intricate app-development framework might.

The most involved part is actually the instruction-display at the top; in order to make it useful and human-readable, having it spit out the instructions as 32-bit binary integers would not cut it. Therefore, the GUI makes use of an additional *display_instruction* function that converts a given machine-code instruction into its equivalent assembly. This feature wasn't too time-intensive, and it was a good use of my growing familiarity with the format of RISC-V's 32-bit instructions. There is also an user-input prompt at the bottom that asks the user to either move to the next step or the previous step, completing the implementation of the rewind feature.

To my surprise, I have found this user display to be rather satisfactory for the project's needs, despite its stark contrast to the initial goal. The only missing feature that I feel is essential, but ran out of time to implement, is an extra line or two for displaying any values that the program writes into the Data Memory. It is imperative to only present modified data rather than the entire storage unit, as otherwise the sheer amount of information would be utterly useless; luckily, the formatting of the Data Memory variable as a hash-map inherently alleviates this concern. It may also be possible to replicate the desired effect of "indicating which wires are passing important information" using color-coded terminal text, but the method used to print the GUI would have made it far too time-consuming to introduce on a whim.

Conclusion

On the whole, this project has been a learning experience in more ways than I ever expected. I've had to start over or rethink my approach to a particular challenge multiple times, and I doubt that I would've grown as familiar with CPU architectures as I have, had I not committed so many mistakes. I would almost recommend that Computer Systems students give their own try at building a CPU simulator rather than just using mine, but I get the feeling that none of them have the extra free month or three to dedicate to such an endeavor. There are a great many lessons and takeaways that I earned in the process of building it, but more on those later. First...

If anyone is, in fact, interested in making use of and building on-top of this simulator, below are some other potential features that could be added as future extensions to the project, even outside of the GUI.

4.1 Future Extensions

The most obvious and tempting features to include are certain ones dealing with ease-of-use: A built-in RISC-V assembler, which would allow users to write and input programs in the much more human-legible format of assembly code; an english "narration" of sorts of what each stage accomplishes on a given cycle; or perhaps even the ability to edit values of the CPU on the fly in order to study how that changes the outcome.

Another approach is to extend the capabilities or fidelity of the simulated CPU. This project limits itself to a simple design and just the RISC-V 32-bit Base Integer Instruction Set for the sake of remaining easy-to-understand for newcomers to Computer Systems, but additional support for some of RISC-V's myriad Standard Extensions could be added to give the simulator some much-needed versatility. These range from things like the inclusion of multiplication instructions and floating-point numbers to vector calculations and a whole 64-bit instruction set, and would come especially handy for anyone studying the RISC-V architecture specifically.

4.1.1 Data Memory

As far as increasing fidelity, the obvious choice is the Data Memory. As described before, the Data Memory in this simulator is a "black box"; it takes in requests to load or store data to certain addresses, and does so without any particular consideration of how that is physically possible. This is technically true for our Register and

Instruction Memories as well, but those are small enough that they can easily fit on the CPU chip, and thus be set up pretty similarly to what our diagram presents. This is not the case for the Data Memory: The Data Memory is magnitudes larger than either of the former two. Even if one only counts the data that a running program access with its 32-bit addresses, the Data Memory can store a mind-boggling 2^{28} words (4-byte chunks of data) compared to the Register Memory's measly 32! That's about 4.2 gigabytes and 128 bytes, respectively; and, since the operating systems of modern personal computers want to be running many, many programs at once, the actual size tends to be even larger in practice, nowadays. Therefore, a possible extension of the simulator would be to alter it to replicate the following more realistic setup: One where the actual Data Memory is stored off the CPU chip, massive but painfully slow to access; while a series of caches mediate this access to hopefully maintain on-chip speed for the more frequently-used parts of Data Memory. This setup would allow the simulator to be usable for learning to work with memory structures as well as CPUs, and could implement crucial features for it such as the aforementioned caches or Virtual Memory.

If this feature is implemented, this will also break another assumption that the current simulator makes: That every pipeline stage takes roughly the same amount of time to execute. It may be worth adding in another subsystem that can measure the time transpired by the CPU's operation in actual milliseconds rather than abstract cycles, to show how significant a time loss total cache misses incur.

4.1.2 Out-of-Order Execution and Register Renaming

If one wishes to extend the CPU with more advanced features, two relatively-simple ones that go hand-in-hand are OOE (Out-of-Order Execution) and Register Renaming.

OOE is simple in concept: Pipelined CPUs sometimes need to stall. In real hardware, CPUs sometimes need to stall a *lot*, especially whenever there's any kind of cache miss in the Data Memory. And many times, multiple of the instructions immediately following one of these costly and stall-heavy instructions (like a Load to an address that's not in cache) are totally unrelated to the slow instruction; like, say, a bunch of arithmetic operations that don't interact with any of the same registers that the costly instruction does. So, why not keep the CPU running while the Data Memory is busy fetching, and get those instructions out of the way? It's out-of-order, but still correct and much faster so long as any instruction that shares registers with the slow instruction stalled out properly the whole way through; these are potential new Data Hazards that OOE brings along with it. This feature should only be implemented after the higher-fidelity Data Memory and its increased time costs, as that's where the biggest performance increase from this technique will come from. If one does this, it may be worth adding an extra piece of data display to note which instruction the Data Memory is diligently working at in the background, while the rest of the CPU continues to run. There are many involved pieces to implementing OOE, but the basic idea is to look at a "window" of the next few instructions and randomly pick one to execute that would not cause a data hazard.

Similarly, Register Renaming should only be implemented after OOE, because Register Renaming is to OOE exactly as Forwarding is to Pipelining: It is a feature that purely exists to sidestep the former's "false" Data Hazards. It goes a little like this: OOE's Data Hazards are much more expansive than Pipelining, due to it changing the ordering of instructions around, divided into three broad categories. Looking back at our earlier example:

1. If the Load instruction loads the value 13 from memory into register $\$r1$, and then an Add-Immediate instruction right after it calculates the value 200 into $\$r1$, the correct behavior would be for the load to finish executing, then for the add-immediate to execute for a final result of $\$r1 = 200$. However, with OOE, the ALU-using instruction would finish executing, and thus finish writing back, way before the Load, and the final value of $\$r1$ would be 13! This is a **WAW** Write-After-Write hazard, since both instructions wrote to the same register.
2. If the Load instruction loads the value 13 from memory into register $\$r1$, and then an Add instruction right after it adds $\$r0$ and $\$r1$ together, one would expect first 13 to be loaded into $\$r1$, and then for the Add to add 0 and 13 together to store in wherever the destination register is. However, with OOE, the Add instruction will instead read whichever value was in $\$r1$ *before* the Load, and produce a totally wrong addition! This is a **RAW**, Read-After-Write hazard.
3. This one requires a full re-ordering, but that is fairly common in OOE. Imagine that the Load instruction loads its data from the address 1024($\$r3$), with $\$r3$ currently holding 0, and that an Add-Immediate instruction after it stores its result, 28 to $\$r3$. It is expected that the Load instruction will load from address 1024(0), but if they are executed out of order, the Add-Immediate will edit $\$r3$ first and the Load instruction will load from 1024(28), a completely wrong address! This is a **WAR** Write-After-Read hazard.

These three sound pretty damning for OOE: instructions depending on each other's registers happens constantly! However, just like in Pipelining, some of these Hazards aren't truly unavoidable. In fact, out of the three, only **RAW** hazards, where the data desired by the Read instruction does not exist until the Write instruction computes, are true! For the WAW and WAR hazards... what if the latter instruction, the Write, (as in, the instruction with a higher Program Count,) simply wrote to a different register entirely, and every instruction afterwards used this new register instead? That would resolve the hazard entirely, without actually changing the outcome of the program! Register Renaming is a feature designed to explicitly enable this, without needing to actually edit the ran program! The general gist is as follows:

In a register-renaming Register Renaming, the $\$r1$ – $\$r31$ general-purpose registers are not actual registers, but rather "names". The actual, physical registers that hold data are called $p1, p2, p3, \dots$. The manufacturer decides how many physical registers there are, but there should absolutely be more than 31. It is important to note that

the program still interacts exclusively with the 32 "name" registers, despite them being just names. Whenever a value is written to a register name, the value actually gets written to an empty physical register, and the register name used is mapped to that physical register in a "mapping table" in the Register Memory; any time a value needs to be read from a register, the Register Memory looks to its mapping table to see which physical register holds the data it needs. Most crucially for OOE, *when a new value is written to a register name that already is mapped to a physical register, that physical register is not modified; instead, the new value is written to a new, empty physical register, and the mapping of that register name is updated to point to the new physical register instead.* Even more crucially, the previous, no-longer-valid mapping remains stored on the table, in the row above. However, this now means that WAW hazard instructions can both store their values without the former instruction incorrectly overwriting the latter one; and it means that the data the Read instruction wants in a WAR hazard is still available!

If you want to implement Register Renaming to this simulator, it may be tempting to try to do so using pointers/references; after all, the register names in this feature are shockingly akin to the references of a programming language. However, this becomes increasingly difficult as one realizes that the Register Memory must keep track of past mappings. I suggest simply making the Mapping Table as an actual table using whatever data structure is most fit for such in the used language; it will make it easier to incorporate the Mapping Table into the GUI as well, which is a must if this feature is implemented.

4.2 Takeaways

Some of the lessons I've acquired are fairly general to managing a project like this, while others are more specific to my particular topic. The biggest such lesson, to me, is that you shouldn't be afraid of ripping initial attempts down and starting over from scratch. It's only natural that the first design you come up with for a program or feature isn't going to be the best, and you can use the experience you've gained in working on the initial attempt to further improve the next design. I've had to do this many, many times throughout the building of this simulator: The window-based GUI was replaced with a command-line interface, the overall code of the program had to be totally redone, and even the finalized design had moments where a function was reordered around to avoid desyncing issues. There are even a few design restarts that I would implement right now, if I had the time: The main one that comes to mind would be reimplementing the Stalling code from scratch, as I'm fairly certain that my choice to implement it outside of the step function is part of what made Load-ALU Stalling so difficult to pull off correctly.

Related, but not exactly identical, is the ability to scale back one's project. My initial design was rather ambitious for what someone might be able to accomplish in a year, especially as I had zero prior experience with either building simulators or working with application development frameworks. However, there's something I partially lucked out on which I came to appreciate later: Many parts of my ambitious

design were modular, such that I could remove or scale down a specific piece of it while still ending up with a complete project. The aforementioned GUI got scaled down to a command-line interface, the detailed Data Memory simulator became a simple black box that merely served the purpose the CPU needed it to, and periphery features like the assembler and wire color-coding were left out entirely. Treating this loss of potential functionality as a good thing may seem counter-intuitive, but setting these aside allowed me to focus on the parts of the project that I really considered essential: The core simulator is implemented and functional save for the one edge-case bug, and the three features I considered critical for the project's purpose (user-inputted programs, back-and-forward stepping, and a graphical display) are all fully functional. Even in more monolithic projects, having aspects of it that can be simplified or removed if needed are incredibly helpful for this. And, it's not like that effort is wasted, as shown by the Extensions section right before this one.

As far as the simulator itself is concerned, I was surprised at which aspects of the CPU design were important to incorporate into a good simulator design, and which ones weren't. The separation of discrete components is actually not that crucial; it is important to keep in mind because the wires and "subcomponents" you wish to depict decide what variables you will have, but actually compartmentalizing the different pieces of CPU codes just causes problems and nightmares. This lines up decently well with the understanding that the separation of "subcomponents" is in itself partly arbitrary, like with how the early CPU diagrams in the background chapter have a singular decoder component, while the simulator splits that into two with the immediates decoder. All that really matters is making sure that you are storing the information you need for the GUI to display, and that the commenting reflects which code snippets belong to each arbitrary piece of the CPU.

On the flipside, the aspect of what parts of the CPU are state-holding and which ones are pure logic is something that I didn't pay much attention to early in the project, but that became massively useful later on! It is an easy pitfall to forget about this since, in the simulated version of this CPU, everything technically has state: We need to record the value coursing through every wire in order to display them in the GUI inbetween steps, and so that is stored in variables. However, it is critical to understand that these wires and stateless subcomponents should still be treated as stateless; It is what makes the final design possible, the idea of updating the state first and then updating all the logic that flows from that.

I also learned that, as far as actually storing this information, simpler is better. Early on, I noted the similarity between subcomponents and the wires connecting them and objects in a programming language with pointers to one another, and thought it would be quite elegant and clever to design my simulator that way; however, we've already gone over how that only caused nightmares in the long run. For future extensions, I recommend writing variables that store the data you need as simply as possible, as I was served so well by simple unsigned 32-bit integers and vectors, and putting them in places where they are easy to access. Making sub-structs for the sake of organization is good, which I ended up doing based on the pipeline stages, but those should all be owned by a small handful of main structs that are easy to feed into functions. For my design, that is simply the State and Logic structs.

There may be more than two main structures in an extended version of this project; for example, a faithfully-simulated Data Memory will certainly be complex enough to warrant distinction from the other stateful components. However, you should consider whether this new struct would be better served by being on its own, or by being a sub-struct to the existing main State struct.

4.3 Closing Thoughts

All in all, I am rather proud of what I've ended up building: A simple, tightly-focused RISC-V simulator that, while lacking in the fancier features of its contemporaries, has enough of a graphical interface to function as an educational tool. The simulator is now freely available at <https://github.com/FakeSalamander/riscv-psim>, and can be compiled and ran on any modern computer with a functioning Rust setup.

There's even something of a hidden silver lining in the project's eventual simplicity: While a simulator like RIPES is meant to be powerful and convenient for the end-user, a more bare-bones commandline program is much easier for an interested undergraduate student to take apart and build on top of. I hope that I have laid down a foundation for what could one day be a polished enough program to potentially see use in a classroom; though even if it simply ends up that a handful of individuals clone its repository to mess around with my code, its purpose will have been accomplished all the same.

References

- Denning, P. J. (1993). The science of computing: Risc architecture. *American Scientist*, 81(1), 7–10. <http://www.jstor.org/stable/29774812>
- Kanter, D. (2013). Silvermont, intel's low power architecture. <https://www.realworldtech.com/silvermont/>
- Kanter, D. (2014). Amd's jaguar microarchitecture. <https://www.realworldtech.com/jaguar/>
- msykpshinz (2019). riscv-isa-pages. <https://msyksphinz-self.github.io/riscv-isadoc/html/rvi.html>
- Price, C. (1995). Mips iv instruction set revision 3.2. <https://www.cs.cmu.edu/afs/cs/academic/class/15740-f97/public/doc/mips-isa.pdf>
- Shimpi, A. L., & Wilson, D. (2004). Intel's pentium 4 e: Prescott arrives with luggage. <https://www.anandtech.com/show/1230/3>
- Waterman, A., Asanovic, K., & Inc., S. (2017). The risc-v instruction set manual volume 1: User-level isa. <https://riscv.org/wp-content/uploads/2017/05/riscv-spec-v2.2.pdf>